

Artificial Intelligence

K. N Toosi University of Technology

Course Instructor:

Dr. Omid Azarkasb

Teaching Assistants:

Atena Najaf Abadi Farahani

Saeed Mahmoudian

Uninformed Search

Search

- Search permeates all of AI
- What choices are we searching through?
 - Problem solving
Action combinations (move 1, then move 3, then move 2...)
 - Natural language
Ways to map words to parts of speech
 - Computer vision
Ways to map features to object model
 - Machine learning
Possible concepts that fit examples seen so far
 - Motion planning
Sequence of moves to reach goal destination
- An intelligent agent is trying to find a set or sequence of actions to achieve a goal
- This is a goal-based agent

Simple Problem Solving Agent

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*) **return** an action

static: *seq* , an action sequence

state, some description of the current world state

goal, a goal

problem, a problem formulation

state \leftarrow UPDATE-STATE(*state*, *percept*)

if *seq* is empty **then**

goal \leftarrow FORMULATE-GOAL(*state*)

problem \leftarrow FORMULATE-PROBLEM(*state* , *goal*)

seq \leftarrow SEARCH(*problem*)

action \leftarrow FIRST(*seq*)

seq \leftarrow REST(*seq*)

return *action*

Problem-solving Agent

```
SimpleProblemSolvingAgent(percept)
  state = UpdateState(state, percept)
  if sequence is empty then
    goal = FormulateGoal(state)
    problem = FormulateProblem(state, g)
    sequence = Search(problem)
  action = First(sequence)
  sequence = Rest(sequence)
  Return action
```

Assumptions

- Static or dynamic?

Environment is static

Assumptions

- Static or dynamic?
- Fully or partially observable?

Environment is fully observable

Assumptions

- Static or dynamic?
- Fully or partially observable?
- Discrete or continuous?

Environment is discrete

Assumptions

- Static or dynamic?
- Fully or partially observable?
- Discrete or continuous?
- Deterministic or stochastic?

Environment is deterministic

Assumptions

- Static or dynamic?
- Fully or partially observable?
- Discrete or continuous?
- Deterministic or stochastic?
- Episodic or sequential?

Environment is sequential

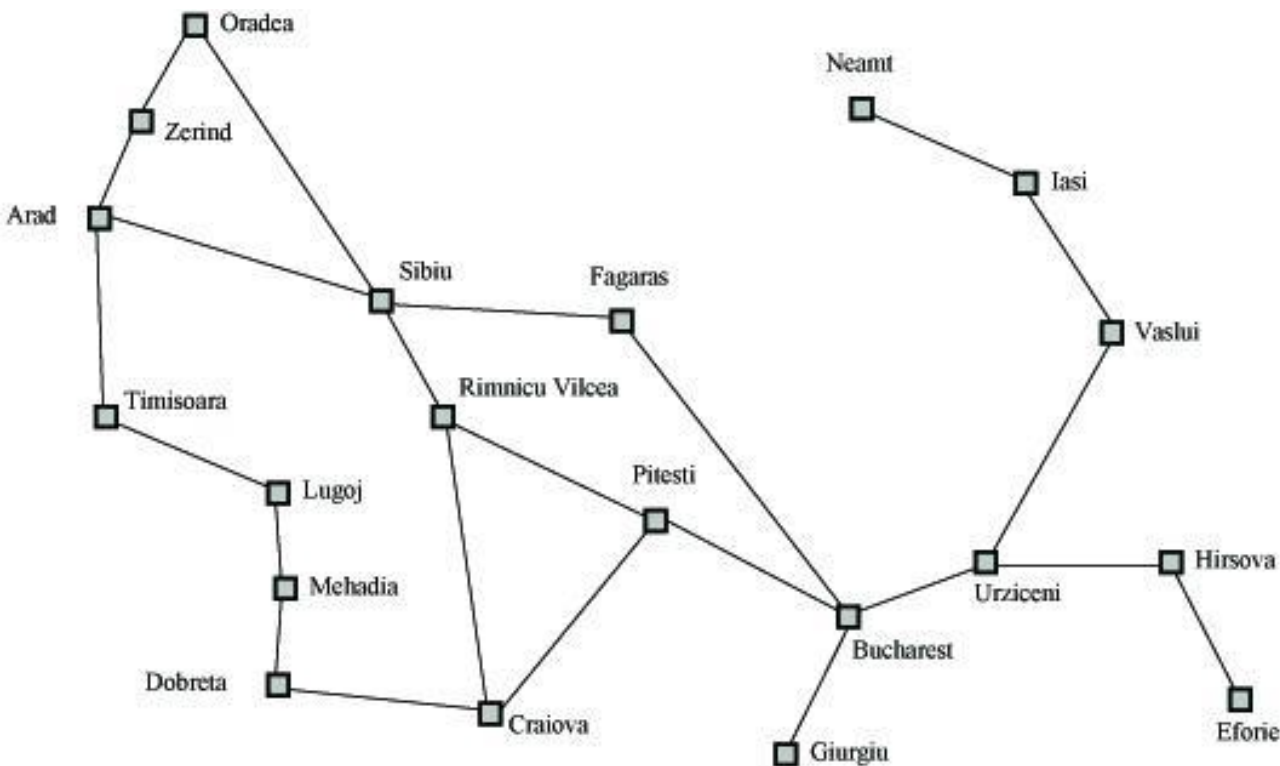
Assumptions

- Static or dynamic?
- Fully or partially observable?
- Discrete or continuous?
- Deterministic or stochastic?
- Episodic or sequential?
- Single agent or multiple agent?

Assumptions

- Static or dynamic?
- Fully or partially observable?
- Discrete or continuous?
- Deterministic or stochastic?
- Episodic or sequential?
- Single agent or multiple agent?

Search Example



Formulate goal: Be in Bucharest.

Formulate problem: states are cities, operators drive between pairs of cities

Find solution: Find a sequence of cities (e.g., Arad, Sibiu, Fagaras, Bucharest) that leads from the current state to a state meeting the goal condition

Search Space Definitions

- **State**
 - A description of a possible state of the world
 - Includes all features of the world that are pertinent to the problem
- **Initial state**
 - Description of all pertinent aspects of the state in which the agent starts the search
- **Goal test**
 - Conditions the agent is trying to meet (e.g., have \$1M)
- **Goal state**
 - Any state which meets the goal condition
 - Thursday, have \$1M, live in NYC
 - Friday, have \$1M, live in Valparaiso
- **Action**
 - Function that maps (transitions) from one state to another

Search Space Definitions

- Problem formulation
 - Describe a general problem as a search problem
- Solution
 - Sequence of actions that transitions the world from the initial state to a goal state
- Solution cost (additive)
 - Sum of the cost of operators
 - Alternative: sum of distances, number of steps, etc.
- Search
 - Process of looking for a solution
 - Search algorithm takes problem as input and returns solution
 - We are searching through a space of possible states
- Execution
 - Process of executing sequence of actions (solution)

Problem Formulation

A search problem is defined by the

1. Initial state (e.g., Arad)
2. Operators (e.g., Arad -> Zerind, Arad -> Sibiu, etc.)
3. Goal test (e.g., at Bucharest)
4. Solution cost (e.g., path cost)

Example Problems – Eight Puzzle

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

States: tile locations

Initial state: one specific tile configuration


Operators: move blank tile left, right, up, or down

Goal: tiles are numbered from one to eight around the square

Path cost: cost of 1 per move (solution cost same as number of most or path length)

[Eight puzzle applet](#)

Example Problems – Robot Assembly



States: real-valued coordinates of


- robot joint angles
- parts of the object to be assembled

Operators: rotation of joint angles

Goal test: complete assembly

Path cost: time to complete assembly

Example Problems – Towers of Hanoi



States: combinations of poles and disks

Operators: move disk x from pole y to pole z
subject to constraints

- cannot move disk on top of smaller disk
- cannot move disk if other disks on top

Goal test: disks from largest (at bottom) to smallest on goal pole

Path cost: 1 per move

[Towers of Hanoi applet](#)

Example Problems – Rubik's Cube



States: list of colors for each cell on each face

Initial state: one specific cube configuration


Operators: rotate row x or column y on face z direction a

Goal: configuration has only one color on each face

Path cost: 1 per move

[Rubik's cube applet](#)

Example Problems – Eight Queens



States: locations of 8 queens on chess board

Initial state: one specific queens configuration

Operators: move queen x to row y and column z

Goal: no queen can attack another (cannot be in same row, column, or diagonal)

Path cost: 0 per move

[Eight queens applet](#)

Example Problems – Missionaries and Cannibals



States: number of missionaries, cannibals, and boat on near river bank

Initial state: all objects on near river bank

Operators: move boat with x missionaries and y cannibals to other side of river


- no more cannibals than missionaries on either river bank or in boat
- boat holds at most m occupants

Goal: all objects on far river bank

Path cost: 1 per river crossing

[Missionaries and cannibals applet](#)

Example Problems –Water Jug



States: Contents of 4-gallon jug and 3-gallon jug

Initial state: (0,0)

Operators:

- fill jug x from faucet
- pour contents of jug x in jug y until y full
- dump contents of jug x down drain

Goal: (2,n)

Path cost: 1 per fill

[Saving the world, Part I](#)


[Saving the world, Part II](#)

Sample Search Problems

- Graph coloring
- Protein folding
- Game playing
- Airline travel
- Proving algebraic equalities
- Robot motion planning

Visualize Search Space as a Tree

- States are nodes
- Actions are edges
- Initial state is root
- Solution is path from root to goal node
- Edges sometimes have associated costs
- States resulting from operator are children




Search Problem Example (as a tree)


(a) The initial state

Arad

(b) After expanding Arad



(c) After expanding Sibiu



Search Function – Uninformed Searches

```
Open = initial state                                // open list is all generated states
                                                    // that have not been “expanded”
While open not empty                                // one iteration of search algorithm
    state = First(open)                             // current state is first state in open
    Pop(open)                                       // remove new current state from open
    if Goal(state)                                 // test current state for goal condition
        return “succeed”                          // search is complete
                                                    // else expand the current state by
                                                    // generating children and
                                                    // reorder open list per search strategy
    else open = QueueFunction(open, Expand(state))
Return “fail”
```

Tree Search Algorithm

```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    else choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  end do
```

Tree Search Algorithm

```
function TREE-SEARCH(problem, fringe) return a solution or failure
  fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if EMPTY?(fringe) then return failure
    node  $\leftarrow$  REMOVE-FIRST(fringe)
    if      GOAL-TEST[node] then
      return SOLUTION(node)
    else
      fringe  $\leftarrow$  INSERT-ALL(EXPAND(node, problem), fringe)
```





Search Strategies

- Search strategies differ only in QueuingFunction
- Features by which to compare search strategies
 - Completeness (always find solution)
 - Cost of search (time and space)
 - Cost of solution, optimal solution
 - Make use of knowledge of the domain
 - “uninformed search” vs. “informed search”

Breadth-First Search

- Generate children of a state, QueueingFn adds the children to the **end** of the open list
- Level-by-level search
- Order in which children are inserted on open list is arbitrary
- In tree, assume children are considered left-to-right unless specified differently
- Number of children is “branching factor” b

BFS Examples



$b = 2$

[Example trees](#)

[Search algorithms applet](#)

Analysis

- Assume goal node at level d with constant branching factor b
- Time complexity (measured in #nodes generated)
 - 1 (1st level) + b (2nd level) + b^2 (3rd level) + ... + b^d (goal level) + $(b^{d+1} - b)$
 $= O(b^{d+1})$
- This assumes goal on far right of level
- Space complexity
 - At most majority of nodes at level d + majority of nodes at level $d+1 = O(b^{d+1})$
 - Exponential time and space
- Features
 - Simple to implement
 - Complete
 - Finds shortest solution (not necessarily least-cost unless all operators have equal cost)

Analysis

- See what happens with $b=10$
 - expand 10,000 nodes/second
 - 1,000 bytes/node


Depth	Nodes	Time	Memory
2	1110	.11 seconds	1 megabyte
4	111,100	11 seconds	106 megabytes
6	10^7	19 minutes	10 gigabytes
8	10^9	31 hours	1 terabyte
10	10^{11}	129 days	101 terabytes
12	10^{13}	35 years	10 petabytes
15	10^{15}	3,523 years	1 exabyte

Depth-First Search

- QueueingFn adds the children to the **front** of the open list
- BFS emulates FIFO queue
- DFS emulates LIFO stack
- Net effect
 - Follow leftmost path to bottom, then backtrack
 - Expand deepest node first

DFS Examples

[Example trees](#)



Analysis

- Time complexity
 - In the worst case, search entire space
 - Goal may be at level d but tree may continue to level m , $m \geq d$
 - $O(b^m)$
 - Particularly bad if tree is infinitely deep
- Space complexity
 - Only need to save one set of children at each level
 - $1 + b + b + \dots + b$ (m levels total) = $O(bm)$
 - For previous example, DFS requires 118kb instead of 10 petabytes for $d=12$ (10 billion times less)
- Benefits
 - May not always find solution
 - Solution is not necessarily shortest or least cost
 - If many solutions, may find one quickly (quickly moves to depth d)
 - Simple to implement
 - Space often bigger constraint, so more usable than BFS for large problems

Comparison of Search Techniques

	DFS	BFS
Complete	N	Y
Optimal	N	N
Heuristic	N	N
Time	b^m	b^{d+1}
Space	bm	b^{d+1}




Avoiding Repeated States

Can we do it?

- Do not return to parent or grandparent state
 - In 8 puzzle, do not move up right after down
- Do not create solution paths with cycles
- Do not generate repeated states (need to store and check potentially large number of states)

Maze Example

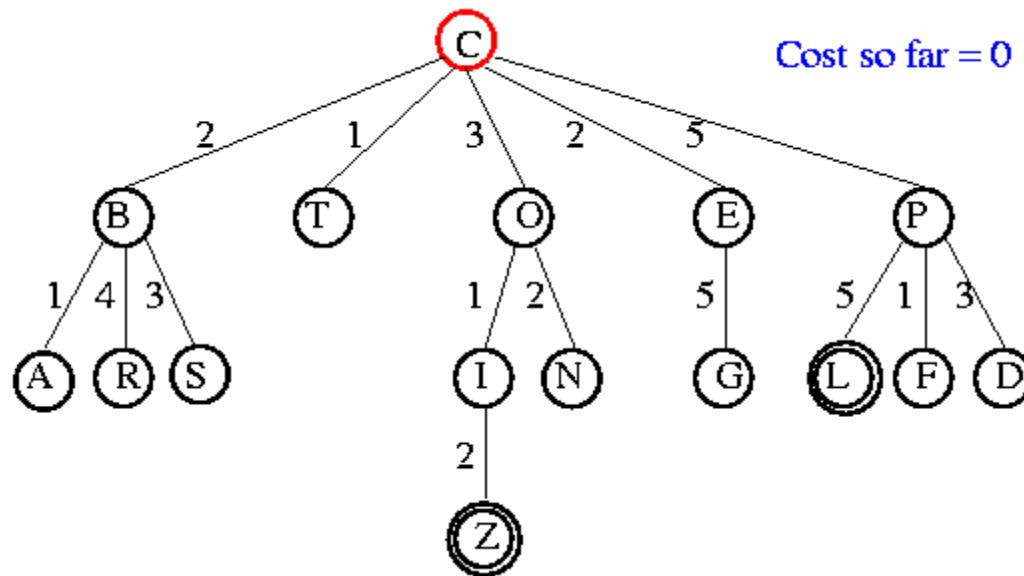
- States are cells in a maze
- Move N, E, S, or W
- What would BFS do (expand E, then N, W, S)?
- What would DFS do?
- What if order changed to N, E, S, W and loops are prevented?



Uniform Cost Search (Branch&Bound)


- QueueingFn is SortByCostSoFar
- Cost from root to current node n is $g(n)$
 - Add operator costs along path
- First goal found is least-cost solution
- Space & time can be exponential because large subtrees with inexpensive steps may be explored before useful paths with costly steps
- If costs are equal, time and space are $O(b^d)$
 - Otherwise, complexity related to cost of optimal solution

UCS Example




Open list: C

UCS Example



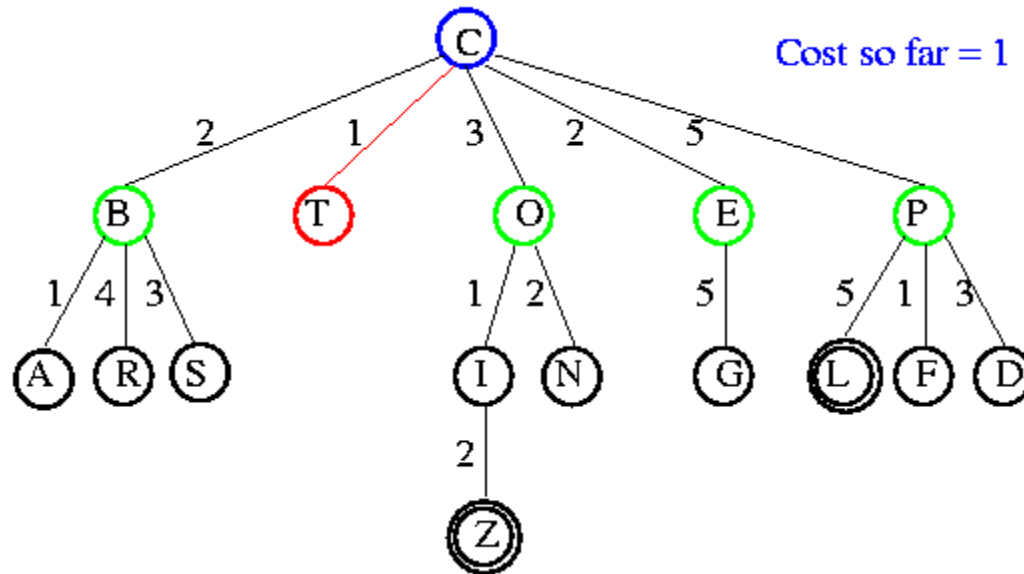
Open list: B(2) T(1) O(3) E(2) P(5)

UCS Example



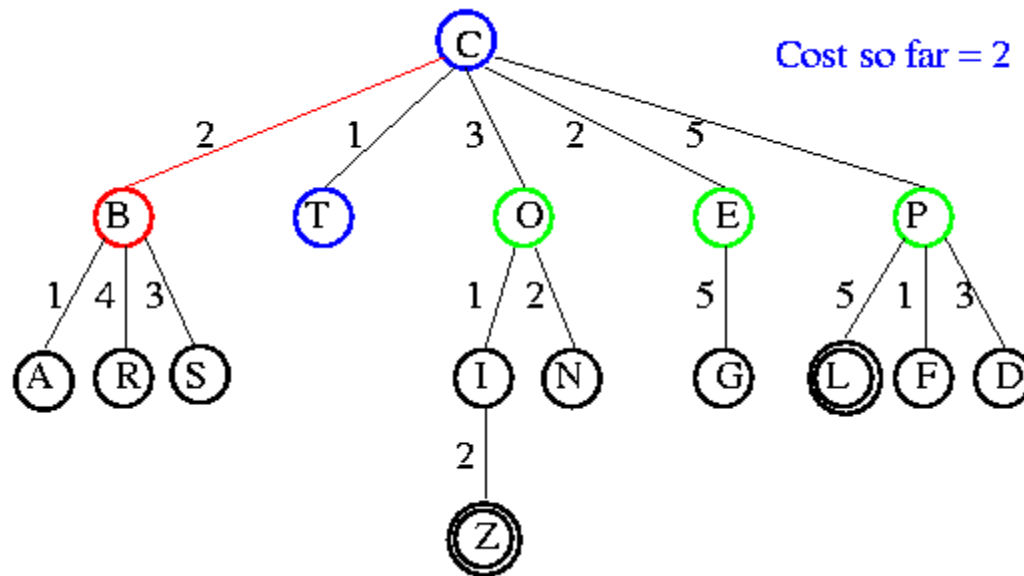
Open list: T(1) B(2) E(2) O(3) P(5)

UCS Example



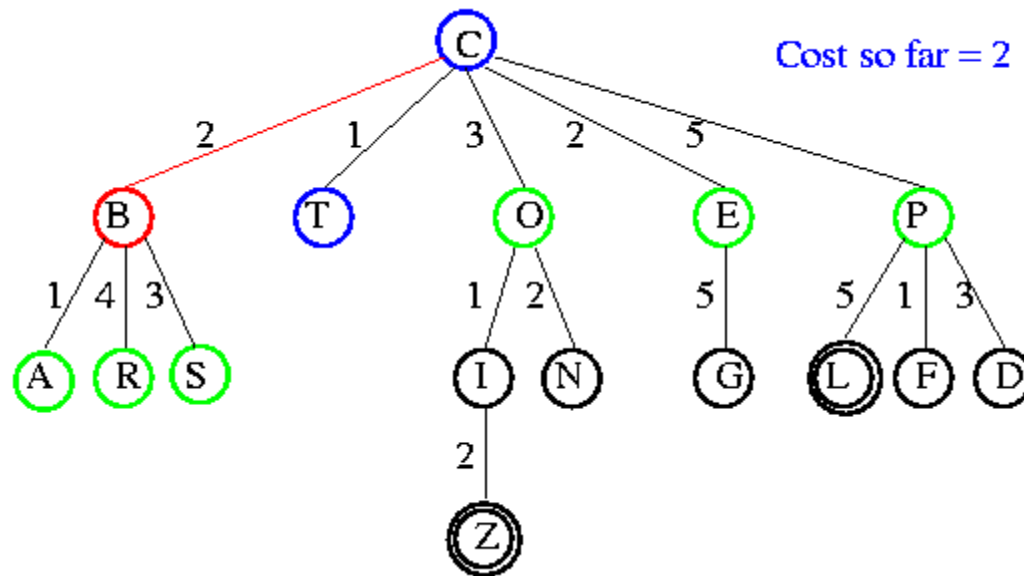
Open list: B(2) E(2) O(3) P(5)

UCS Example



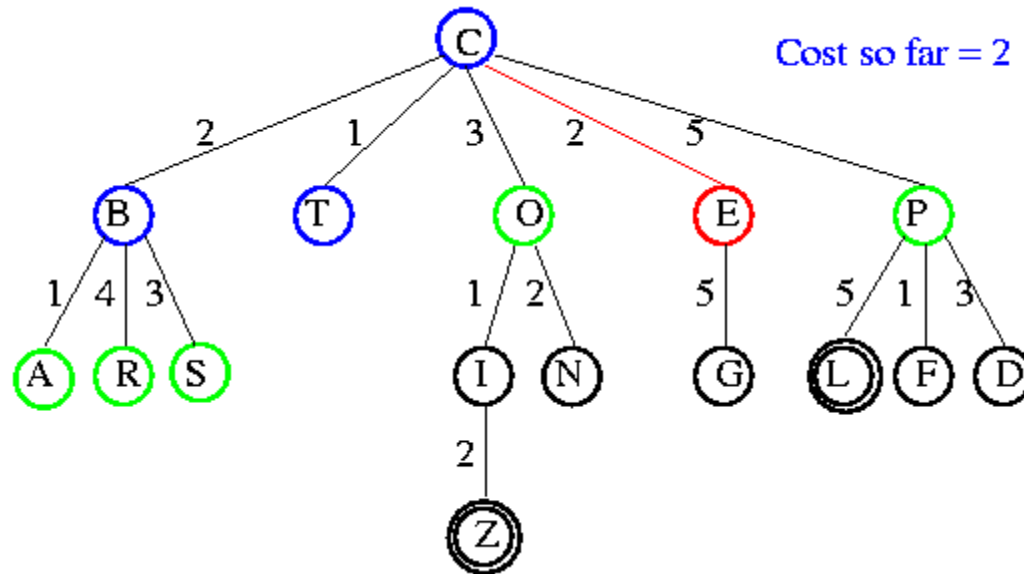
Open list: E(2) O(3) P(5)

UCS Example




Open list: E(2) O(3) A(3) S(5) P(5) R(6)

UCS Example



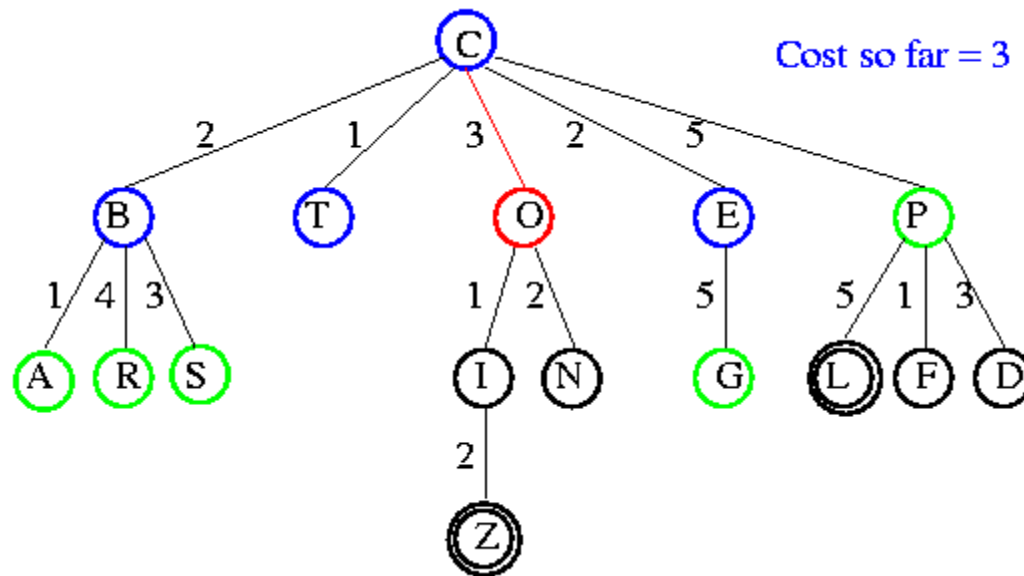
Open list: O(3) A(3) S(5) P(5) R(6)

UCS Example




Open list: O(3) A(3) S(5) P(5) R(6) G(10)

UCS Example



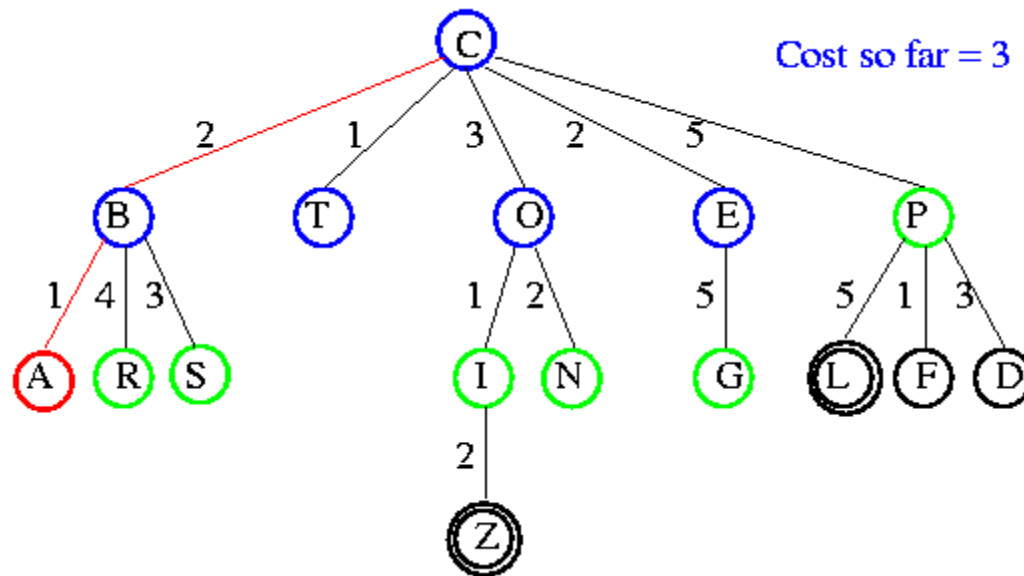
Open list: A(3) S(5) P(5) R(6) G(10)

UCS Example



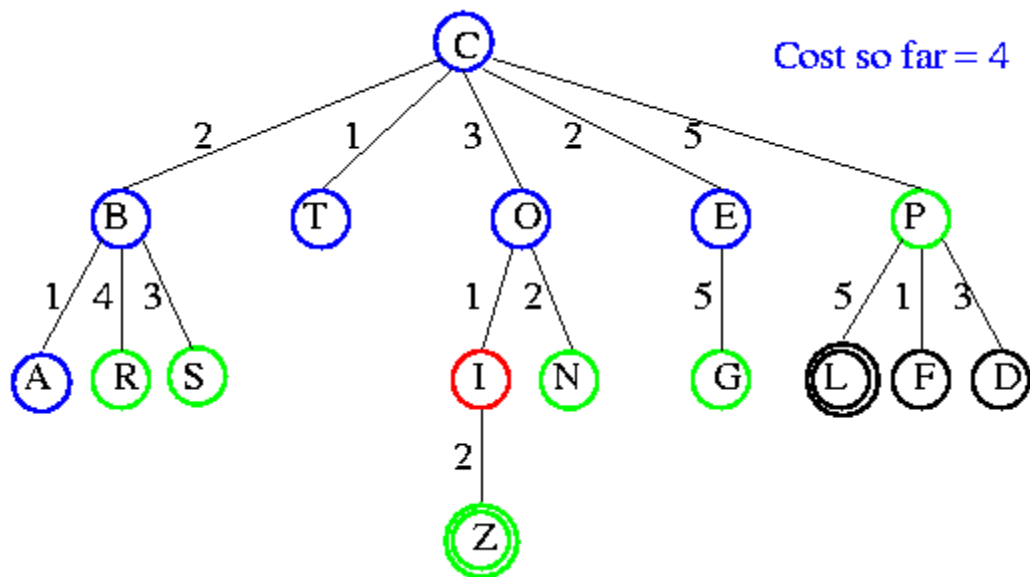
Open list: A(3) I(4) S(5) N(5) P(5) R(6) G(10)

UCS Example



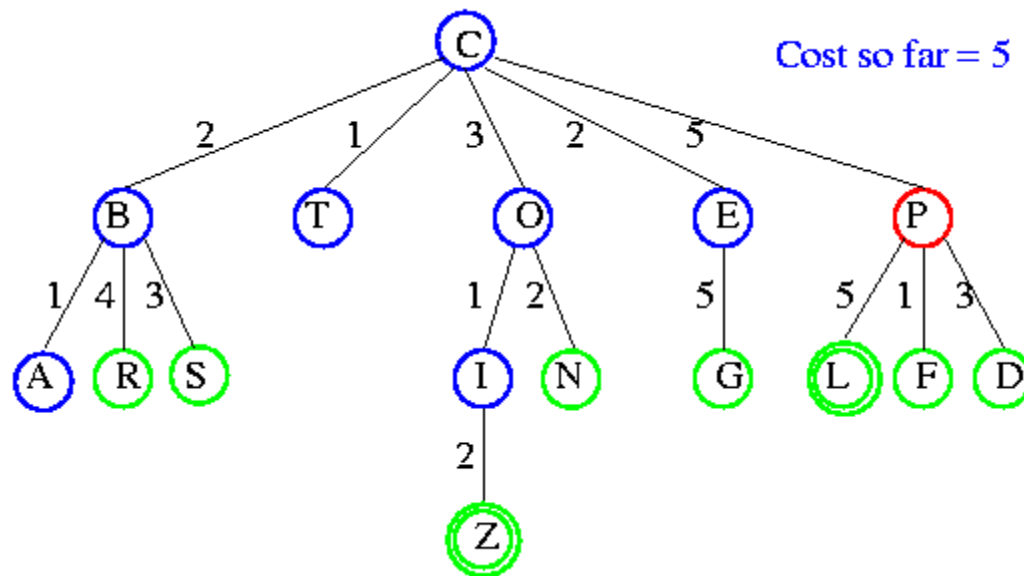
Open list: I(4) P(5) S(5) N(5) R(6) G(10)

UCS Example



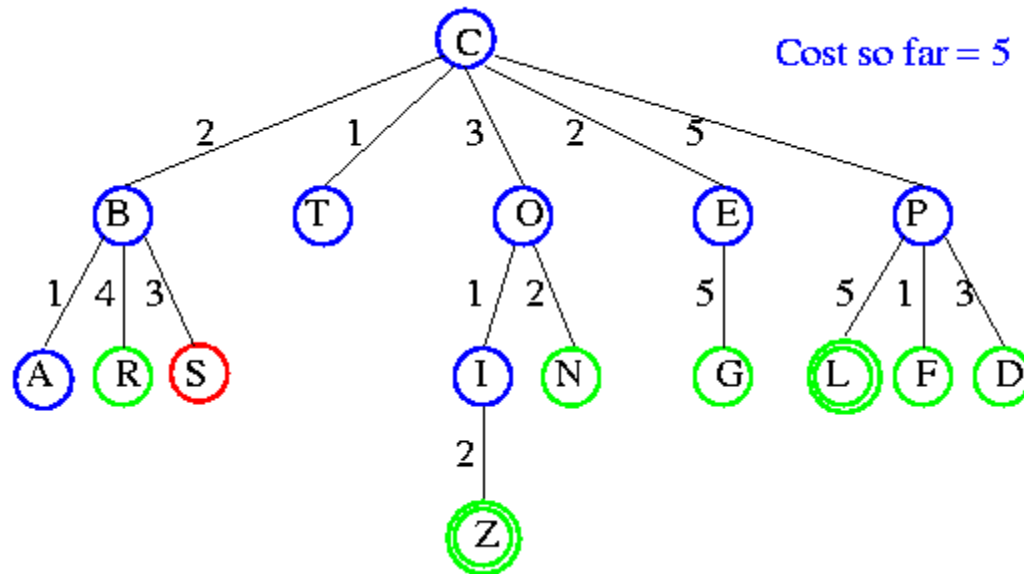
Open list: P(5) S(5) N(5) R(6) Z(6) G(10)

UCS Example



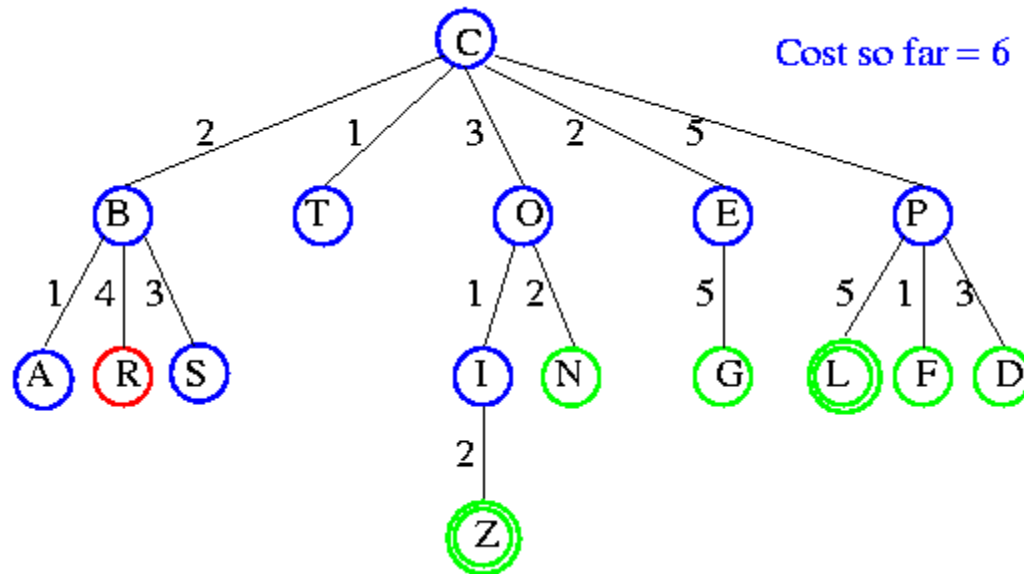
Open list: S(5) N(5) R(6) Z(6) F(6) D(8) G(10) L(10)

UCS Example




Open list: N(5) R(6) Z(6) F(6) D(8) G(10) L(10)

UCS Example



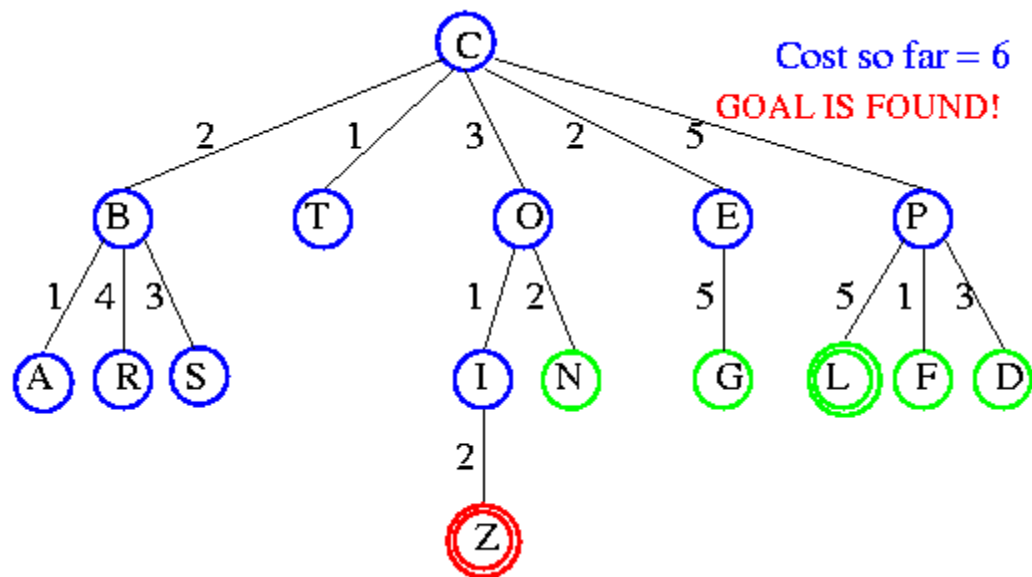
Open list: Z(6) F(6) D(8) G(10) L(10)

UCS Example



Open list: F(6) D(8) G(10) L(10)

UCS Example



Comparison of Search Techniques

	DFS	BFS	UCS
Complete	N	Y	Y
Optimal	N	N	Y
Heuristic	N	N	N
Time	b^m	b^{d+1}	b^m
Space	bm	b^{d+1}	b^m



IDS

(Iterative Deepening Search)




- DFS with depth bound
- QueuingFn is enqueue at front as with DFS
 - Expand(state) only returns children such that $\text{depth}(\text{child}) \leq \text{threshold}$
 - This prevents search from going down infinite path
- First threshold is 1
 - If do not find solution, increment threshold and repeat

Examples









Limit = 0















Limit = 1



Limit = 2



Limit = 3



Analysis

- What about the repeated work?
- Time complexity (number of generated nodes)
 - $[b] + [b + b^2] + \dots + [b + b^2 + \dots + b^d]$
 - $(d)b + (d-1) b^2 + \dots + (1) b^d$
 - $O(b^d)$

Analysis

- Repeated work is approximately $1/b$ of total work
 - Negligible
 - Example: $b=10$, $d=5$
 - $N(\text{BFS}) = 1,111,100$
 - $N(\text{IDS}) = 123,450$
- Features
 - Shortest solution, not necessarily least cost
 - Is there a better way to decide threshold? (IDA^*)


Comparison of Search Techniques

	DFS	BFS	UCS	IDS
Complete	N	Y	Y	Y
Optimal	N	N	Y	N
Heuristic	N	N	N	N
Time	b^m	b^{d+1}	b^m	b^d
Space	bm	b^{d+1}	b^m	bd



Bidirectional Search

- Search forward from initial state to goal AND backward from goal state to initial state
- Can prune many options
- Considerations
 - Which goal state(s) to use
 - How determine when searches overlap
 - Which search to use for each direction
 - Here, two BFS searches
- Time and space is $O(b^{d/2})$



Uninformed Search Algorithms Comparison

Criterion	Breadth-First	Uniform-cost	Depth-First	Depth-limited	Iterative deepening	Bidirectional search
Complete?	YES*	YES*	NO	YES, if $l \geq d$	YES	YES*
Time	b^{d+1}	$b^{C^*/\epsilon}$	b^m	b^l	b^d	$b^{d/2}$
Space	b^{d+1}	$b^{C^*/\epsilon}$	bm	bl	bd	$b^{d/2}$
Optimal?	YES*	YES*	NO	NO	YES	YES